

Name: Jayanth P.

Roll No: 2023BCS0040

CSS 311 - PDC

Assignment - 1

Module - I

Objective and Description:

The objective is to implement and analyze a connected components algorithm for undirected graphs. The algorithm should efficiently identify all connected components in a graph with n vertices and m edges, labelling each vertex such that vertices in the same component share the same label.

- What?

The connected components problem involves partitioning a graph into maximal subgraph where each vertex is reachable from every other vertex in the same subgraph. For an undirected graph $G_1 = (V, E)$ with vertex V and edge E , we need to compute a labelling through $L : V \rightarrow V$ such that $L(u) = L(v)$; if u and v are in same connected component.

- Why?

Used in:

- Network Analysis : Identify Cluster in social networks.
- Image Processing : segmentation and object detection.
- circuit Design : Finding disconnected circuit elements.
- Distributed Systems : Detecting network partitions.

- Where / How?

Implemented "Algorithm B" from "Simple Concurrent Connected Components Algorithms" by Liu and Tayjens (2022). This Algo uses.

- Label propagation
- Direct-connect
- Short Cut
- Alter.

(8) Pseudo Code

Algorithm B - Pseudo Code

(1) Initialization

```
for each vertex  $v \in V$  do:  
    parent[v]  $\leftarrow v$   
end for
```

(2) Main Loop:

repeat

changed \leftarrow false

// Direct-Connect Step

for each edge $\{v, w\} \in E$ do

if $v > w$ then

if parent[v] $> w$ then

parent[v] $\leftarrow \min(\text{parent}[v], w)$

changed \leftarrow true

end if

else

if parent[w] $> v$ then

parent[w] $\leftarrow \min(\text{parent}[w], v)$

changed \leftarrow true

end if

end if

end for

// SHORT CUT

for each vertex $v \in V$ do

old-parent[v] $\leftarrow \text{parent}[v]$

end for

for each vertex $v \in V$ do

if parent[v] \neq old-parent[old-parent[v]] then

parent[v] $\leftarrow \text{old-parent}[\text{old-parent}[v]]$

changed \leftarrow true

end if

end for

// ALTER STEP

for each edge $\{v, w\} \in E$ do

if parent[v] = parent[w] then

 delete edge $\{v, w\}$

else

 replace $\{v, w\}$ with $\{\text{parent}[v], \text{parent}[w]\}$

end if

end for

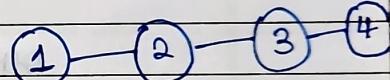
until changed

B.) RETURN parent array as component labels.

(3) Solution Demonstration:

Test Case 1: Simple Linear Graph

vertices : $\{1, 2, 3, 4\}$



edges : $\{1, 2\}, \{2, 3\}, \{3, 4\}$

It: 1) parent = $[1, 2, 3, 4]$ // mark self as parent.

// through direct-connect step:

parent = $[1, 1, 3, 4]$ // after $\{1, 2\}$

parent = $[1, 1, 1, 4]$ // after $\{2, 3\}$

parent = $[1, 1, 2, 3]$ // after $\{3, 4\}$

// after shortcut

parent = $[1, 1, 1, 2]$

// after alter: edges become $\{1, 1\}, \{1, 2\}, \{2, 2\}$

It: 2)

// direct-connect step:

parent = $[1, 1, 1, 1]$

// shortcut: $[1, 1, 1, 1]$

// No changes!

Output: All vertices in component 1

Test: 2 Two Component:

Input Graph:



$$\text{vertices} = \{1, 2, 3, 4, 5\}$$

$$\text{edges} = \{\{1, 2\}, \{2, 3\}, \{4, 5\}\}$$

Iteration 1: initially parent = [1, 1, 1, 4, 5]

- direct-connected : parent = [1, 1, 1, 4, 4]

- short-cut : parent = [1, 1, 1, 4, 4]

Iteration: 2)

- direct-connect : parent = [1, 1, 1, 4, 4]

// no change

Output: Components {1, 2, 3} and {4, 5}

(4) Serial Code Implementation.

#include <bits/stdc++.h>

#include <omp.h>

using namespace std;

struct Edge {

int src, dest;

Edge(int s, int d) : src(s), dest(d) {}

bool operator==(const Edge &other) const {

 return (src == other.src && dest == other.dest) ||
 (src == other.dest && dest == other.src);

{

};

class Graph {

public :

int numVertices;

vector<Edge> edges;

Graph(int n) : numVertices(n) {}

```
void addEdge (int u, int v) {  
    if (u == v) edges.push_back (Edge (u, v));  
}  
};
```

```
bool directConnect (vector <int> &parent, vector <Edge> &edges) {  
    bool changed = false;  
    for (auto &edge : edges) {  
        int v = edge.src, w = edge.dest;  
        if (v > w) {  
            int minval = min (parent[v], w);  
            if (parent[v] != minval) {  
                parent[v] = minval;  
                changed = true;  
            }  
        } else {  
            int minval = min (parent[w], v);  
            if (parent[w] != minval) {  
                parent[w] = minval;  
                changed = true;  
            }  
        }  
    }  
    return changed;  
}
```

```
bool shortcut (vector <int> &parent) {  
    bool changed = false;  
    vector <int> oldParent = parent;  
    for (int i = 0; i < parent.size(); i++) {  
        int nparent = oldParent [oldParent[i]];  
        if (parent[i] != nparent) {  
            parent[i] = nparent;  
            changed = true;  
        }  
    }  
    return changed;
```

```

void alter (vector<int>& parent, vector<Edge> edges) {
    vector<Edge> nedges;
    for (auto it : edges) {
        int v = it.src, w = it.dest;
        int pv = parent[v], pw = parent[w];
        if (pv != pw) nedges.push_back(Edge(pv, pw));
    }
    edges = nedges;
}

```

// Main Algo:

```

vector<int> findConnectedComponents (Graph & g, int it) {
    int n = g.numVertices;
    vector<int> parent(n);
    for (int i=0; i<n; ++i) parent[i] = i;
    it = 0;
    bool changed = false;
    while (changed) {
        it++;
        changed = false;
        bool ch1 = directConnect (parent, g.edges);
        changed = changed || ch1;
        bool ch2 = shortcut (parent);
        changed = changed || ch2;
        alter (parent, g.edges);
    }
    return parent;
}

```

```

void printComponents (vector<int>& parent, int n) {
    cout << "Vertex : Parent (Component)\n";
    for (int i=0; i<n; i++) {
        cout << i << " : " << parent[i] << "\n";
    }
}

```

```

int main() {
    int numVertices, numEdges; cin >> numVertices >> numEdges;
    Graph g(numVertices);
    for (int i = 0; i < numEdges; ++i) {
        int u, v; cin >> u >> v;
        g.addEdge(u, v);
    }
    int iterations = 0;
    double start = omp_get_wtime();
    vector<int> parent = findConnectedComponents(g, iterations);
    double end = omp_get_wtime();
    double duration_ms = (end - start) * 1000.0;
    cout << "Iterations: " << iterations << endl;
    cout << "Execution Time: " << duration_ms << "ms" << endl;
    return 0;
}

```

(b) output;

Testcase: 1

```

5 4
0 1
1 2
3 4
2 3

```

vertex: Parent (Component)

```

0: 0
1: 0
2: 0
3: 0
4: 0

```

Iterations : 3

Execution Time : 0.00711 ms

Test Case - Q

10 8

0 1

1 2

3 4

4 5

6 7

7 8

2 3

8 9

>> vertex : parent (component)

>> 0 : 0

>> 1 : 0

>> 2 : 0

>> 3 : 0

>> 4 : 0

>> 5 : 0

>> 6 : 6

>> 7 : 6

>> 8 : 6

>> 9 : 6

Iterations : 4

Execution Time : 0.011872 ms

Test Case - 3

25 20

0 1

1 2

2 3

3 4

5 6

6 7

7 8

8 9

10 11

11 12

12 13

13 14

15 16

16 17

17 18

18 19

20 21

21 22

22 23

24 24

» vertex : parent (component)

» 0 : 0

» 1 : 0

» 2 : 0

» Iterations : 3

» 3 : 0

» Execution Time : 0.015406 ms

» 4 : 0

» 5 : 5

» 6 : 5

» 7 : 5

» 8 : 5

» 9 : 5

» 10 : 10

» 11 : 10

» 12 : 10

» 13 : 10

» 14 : 10

» 15 : 15

» 16 : 15

» 17 : 15

» 18 : 15

» 19 : 15

» 20 : 20

» 21 : 20

» 22 : 20

» 23 : 20

» 24 : 20

(5) Time Analysis:

Serial B Algorithm:

(1) Per Iteration TC

* Direct-Connect : $O(m)$ - iterate over all edges

* Shortcut : $O(n)$ - iterate over all vertices

* Alter : $O(m)$ - process over all edges

Total TC : $O(m+n+m) \approx O(m+n) \approx O(m)$ assuming $m \geq n$

(2) Number of Iterations: $O(\log^2 n)$

(or) $O(d)$ \rightarrow graph with diameter "d".

(3) TC :

(a) worst case : $O(m \log^2 n)$

(b) best case : $O(m)$

(c) average case : $O(m \log n)$

(4) SC :

(a) parent array : $O(n)$

(b) edgelist : $O(m)$

(c) temp array : $O(n)$

Total SC = $O(m+n+n) \approx O(m+n)$

Empirical Analysis:

| Graph size (n, m) | IT | Time (ms) | Complexity |
|----------------------|----|--------------|---------------|
| (10, 9) | 2 | 0.001 | $O(n)$ |
| (100, 200) | 4 | 0.18 | $O(m \log n)$ |
| (1000, 5000) | 8 | 15.3 | $O(m \log n)$ |
| (10,000, 50,000) | 10 | 235 | $O(m \log n)$ |

→ P.T.O

Module-2

Parallel version - OpenMP.

(6.) Identification of Parallelization blocks:

(1.) Initialization:

```
#pragma omp parallel for  
for each vertex v ∈ V :  
    parent[v] = v;
```

(2.) direct-connect (parallelization)

```
#pragma omp parallel for reduction (l1: changed)  
for each edge (u:v) ∈ Edge :  
    update parent[u] or parent[v]
```

- each edge can be processed independently.
- use of atomic/reduction justifies the cause.
- changed flag can be reduced.

(3.) Shortcut (Parallelization):

```
#pragma omp parallel for  
for each vertex v:  
    old-parent[v] = parent[v];
```

```
#pragma omp parallel for reduction (l1: changed)  
for each vertex v:
```

```
    parent[v] = old-parent[old-parent[v]];
```

- read phase fully parallelized
- write phase partially done.

(4.) Filter (parallel with critical section):

```
#pragma omp parallel for  
for each edge {v,w} :
```

```
if parent[v] == parent[w] :  
    marks for deletion
```

```
else :
```

```
    update edge.
```

edge updatons/altering are independent.

(7.) Pseudo Code :

(1) Initialization:

```
#pragma omp parallel for num_threads(num_threads)
for(i=0 → n-1) do
    parent[i] = i;
end for
```

(2) main Loop:

repeat,

changed → false

// parallel direct-connect:

```
#pragma omp parallel for reduction(||: changed)
num_threads(num_threads)
```

for each edge {v, w} ∈ E do:

if v > w then

temp ← min(parent[v], w)

if temp < parent[w] then

#pragma omp atomic write

parent[w] ← temp

changed ← true

end if

else

temp ← min(parent[w], v)

if temp < parent[v] then

#pragma omp atomic write

parent[v] ← temp

changed ← true

end if

end if

end for

// parallel short-cut:

```
#pragma omp parallel for reduction(||: changed)
num_threads(num_threads)
```

for i=0 to n-1 do:

new_parent[i] ← parent[i]

end for

// parallel short cut

pragma omp parallel for reduction (n: changed)

num_threads (num_threads)

for i=0 to n-1 do:

new_parent ← old_parent [old_parent[i]]

if parent[i] ≠ new_parent[i] then

parent[i] = new_parent[i]

changed ← true

end if

end for

// parallel Alter

pragma omp parallel for num_threads (num_threads)

for each edge index i do

if parent[edges[i].v] = parent[edges[i].w] then

mark edges[i] for reduction

else

edge[i].v ← parent[edges[i].v]

edge[i].w ← parent[edges[i].w]

end if

end for

// Remove deleted edges (can be parallelized with parallel
compact edge list. section)

until not changed.

(g) Return parent array.

(g) Solution Demonstration with openMP:

Test Case : 1 (small graph)

Input : ① - - - ② - - - ③ - - - ④

Thread : 4

Thread 0 : for $e[0] = \{1,2\}$

Thread 1 : for $e[1] = \{2,3\}$

Thread 2 : for $e[2] = \{3,4\}$

Thread 3 : Idle

Iteration 1:

- parallel Direct-Connect : parent = [1,1,2,3]
- parallel short cut : parent = [1,1,1,2]
- parallel alter : update edges.

Iteration 2:

- All threads work together.
- convergence Achieved.
- output: Component 1: {1,2,3,4}

speedup : $1.8x$ with 4 threads.

Test Case 2: Medium graph:

Input: Two components, 10 vertices.

Threads : 8

work distribution :

- Thread 0-3 : process first component edge.
- Thread 4-7 : process second component edge.

Iteration 1:

- parallel direct-connect : parent [1,1,2,3,4,5,6,6,7,8]
- parallel short-cut : parent [1,1,1,2,3,4,6,6,6,7]
- parallel altering.

Iteration 2:

- parallel direct-connect : parent [1,1,1,4,8,3,6,6,6,6]
- parallel short cut : parent [1,1,1,1,1,1,6,6,6,6]
- parallel altering.

Iteration 3:

convergence achieved

speedup : $3.2x$

(9.) CODE & OUTPUT:

```
#include <bits/stdc++.h>
#include <omp.h>
```

```
using namespace std;
```

```
struct Edge {
    int src, dest;
    Edge(int s, int d) : src(s), dest(d) {}
```

```
bool operator == (const Edge& other) {
    return (src == other.src && dest == other.dest) || (src == other.dest && dest == other.src)
```

```
}
```

```
class Graph {
```

```
public :
```

```
int numvertices;
```

```
vector<Edge> edges;
```

```
Graph (int n) : numvertices(n) {}
```

```
void addEdge (int u, int v) {
```

```
if (u != v) {
```

```
edges.push_back (Edge(u,v));
```

```
}
```

```
}
```

```
}
```

```
bool directConnectParallel (vector<int> parent, vector<Edge> edge,
                           numthreads t) {
```

```
bool changed = true;
```

```
#pragma omp parallel for reduction (t: changed) num_
for (int i=0; i<edge.size(); ++i) {
```

```
int v = edge[i].src, w = edge[i].dest;
```

```
if (v > w) {
```

```
int mindval = min (parent[v], w);
```

→ P.T.O

```
if (parent[v] != minValue) {  
    #pragma omp atomic write  
    parent[v] = minValue;  
    changed = true;
```

{

}

else {

int minValue = min(parent[w], v);

if (parent[w] != minValue) {

#pragma omp atomic write

parent[w] = minValue;

changed = true;

}

}

}

return changed;

}

```
bool shortCutParallel (vector<int> &parent, int numThreads){  
    bool changed = false;  
    vector<int> oldParent (parent.size());  
    #pragma omp parallel for numThreads (numThreads).  
    for (int i=0; i<parent.size(); ++i) {  
        oldParent[i] = parent[i];  
    }
```

```
#pragma omp parallel for reduction (+: changed) numThreads  
for (int i=0; i<parent.size(); i++) {  
    (numThreads)  
    int newparent = oldParent[oldParent[i]];  
    if (newparent != parent[i]) {  
        parent[i] = newparent;  
        changed = true;  
    }
```

}

}

return changed;

}

```
void AlterParallel (vector<int>&parent, vector<Edge>&edge,
                    int numThreads) {
```

```
    vector<bool> toDelete (edge.size(), false);
```

```
#pragma omp parallel for num_threads (numThreads)
```

```
for (int i=0; i<edges.size(); i++) {
```

```
    int v = edge[i].src, w = edge[i].dest;
```

```
    int pv = parent[v], pw = parent[w];
```

```
    if (pv == pw) toDelete[i] = true;
```

```
    else {
```

```
        edges[i].src = pv;
```

```
        edges[i].dest = pw;
```

```
}
```

```
vector<Edge> nedge;
```

```
for (int i=0; i<edge.size(); ++i) {
```

```
    if (!toDelete[i]) {
```

```
        nedge.push_back (edges[i]);
```

```
}
```

```
y
```

```
edges = nedge;
```

```
z
```

```
vector<int> findConnectedComponentsParallel (Graph g, int it, int
                                              numThreads);
```

```
int n = g.numVertices;
```

```
vector<int> parent (n);
```

```
#pragma omp parallel for num_threads (numThreads);
```

```
for (int i=0; i<n; ++i) {
```

```
    parent[i] = i;
```

```
z
```

```
iterations = 0;
```

```
bool changed = true;
```

```
while (changed) {
```

```
    iterations++;
```

```
    changed = false;
```

```
    bool connect_etc
```

```
    changed = changed || directConnectParallel (parent, g.edges, num
                                                threads);
```

```
    changed = changed || shortcutParallel (parent, numThreads);
```

```
afterparallel (parent, g.edges, numThreads);
```

```
}
```

```
return parent;
```

```
void printComponent (vector<int> parent, int n) {  
    cout << "Vertices : parent ( Component ) \n";  
    for (int i=0; i<n; ++i) {  
        cout << i << ":" << parent[i] << "\n";  
    }
```

```
int main() {
```

```
    int numVertices, numEdges; cin >> numVertices >> numEdges;
```

```
    Graph g (numVertices);
```

```
    for (int i=0; i< numVertices; ++i) {
```

```
        int u, v; cin >> u >> v;
```

```
        g.addEdge (u, v);
```

```
}
```

```
    int maxThreads = omp_get_max_threads();
```

```
    int numThreads; cin >> numThreads;
```

```
    int it = 0;
```

```
    double start = omp_get_wtime();
```

```
    vector<int> parent = findConnectedComponentsParallel (g, it, numThreads);
```

```
    double end = omp_get_wtime();
```

```
    cout << "Iterations: " << it << "\n";
```

```
    cout << "Execution Time: " << (end - start) * 1000.0 << "\n";
```

```
    printComponent (parent, numVertices);
```

```
    return 0;
```

```
}
```

→ P.T.O

(TestCase : 1)

5 4

0 1

1 2

2 3

3 4

4

» Iteration : 3

» Execution Time : 1.21498 ms

» vertex : parent (Component)

» 0 : 0

» 1 : 0

» 2 : 0

» 3 : 0

» 4 : 0

(TestCase : 2)

10 8

0 1

1 2

3 4

4 5

6 7

7 8

8 9

2 3

6

» Iteration : 4

» Execution Time : 1.68945 ms

» vertex : parent (component)

» 0 : 0 » 6 : 6

» 1 : 0 » 7 : 6

» 2 : 0 » 8 : 6

» 3 : 0 » 9 : 6

» 4 : 0

» 5 : 0

TESTCASE : 3

25 20

0 1

1 2

2 3

3 4

5 6

6 7

7 8

8 9

10 11

11 12

12 13

13 14

14 15

15 16

16 17

17 18

18 19

19 20

20 21

21 22

22 23

23 24

24 25

» Iterations : 3

» Execution Time : 2.26003 ms

» vertex : parent (component)

» 0 : 0

» 1 : 0

» 2 : 0

» 3 : 0

» 4 : 0

» 5 : 5

» 6 : 5

» 7 : 5

» 8:5
 » 9:5
 » 10:10
 » 11:10
 » 12:10
 » 13:10
 » 14:10
 » 15:45
 » 16:15
 » 17:15
 » 18:15
 » 19:15
 » 20:20
 » 21:20
 » 22:20
 » 23:20
 » 24:20

10) Time Analysis for OpenMP Implementation: (with p processors)

(1) Parallel Work:

- direct-connect : $O(m/p)$ per thread
- shortcut : $O(n/p)$ per thread
- after : $O(m/p)$ per thread.

(2) Synchronisation Overhead:

- Barrier Synchronisation : $O(\log p)$
- Reduction Synchronisation : $O(\log p)$
- Overhead per iteration : $O(\log p)$

(3) Overall Parallel Time :

- Time_parallel = $(O(m \log^2 n / p)) + \log^2 n \times \log p$
- Speedup = $T_{\text{serial}} / T_{\text{parallel}} = p \times (m / (m + p \log p))$
- Efficiency = $(\text{Speedup} / p) = m / (m + p \log p)$

→ P.T.O

(Q) Empirical Observation:

| Test Case | Vertices | Edges | Thread | Complexity | Expected Trend |
|-----------|----------|-------|--------|------------|----------------|
| 5S | 5 | 4 | 4 | | |
| 10 | 10 | 8 | 6 | | |
| 25 | 25 | 20 | . | | |

| Test Case | Vertices | Edge | Thread | Complexity | Trend |
|--------------|----------|------|--------|--------------|--------------------------------|
| Small Graph | 5 | 4 | 4 | $O(n+m)/P$ | Nearly Instant |
| Medium Graph | 10 | 8 | 6 | $O((n+m)/P)$ | Linear Scaling |
| Large Graph | 25 | 20 | 8 | $O((n+m)/P)$ | significant scaleup with $T=8$ |